

Lecture 2

Host programming

Outline

- **Introduction to OpenCL**
- **OpenCL Framework**
 - Platform layer
 - Compiler for OpenCL C
 - Runtime
- **The Flow of Host Program**
 - Example: Hello world

Outline

- Introduction to OpenCL
- **OpenCL Framework**
 - Platform layer
 - Compiler for OpenCL C
 - Runtime
- **The Flow of Host Program**
 - Example: Hello world

What is OpenCL

- **OpenCL (Open Computing Language)**
 - is a framework suited for parallel programming of heterogeneous systems.
- **The framework supports heterogeneous accelerated processing units**
 - execution on CPU, DSPs, FPGAs, GPUs, and other.
- **OpenCL provides parallel computing**
 - task-based parallelism
 - data-based parallelism


Design Goals of OpenCL

- **Use all computational resources in the system**
 - CPUs, GPUs and other processors as peers
- **Efficient parallel programming model**
 - Based on C99
 - Data- and task-parallel computational model
 - Abstract the specifics of underlying hardware
 - Specify accuracy of floating-point computations
- **Desktop and Handheld Profiles**

CPU-GPU Cooperation (1/2)


• CPU

- Optimized for low-latency access to cached data sets
- Control logic for out-of-order and speculative execution



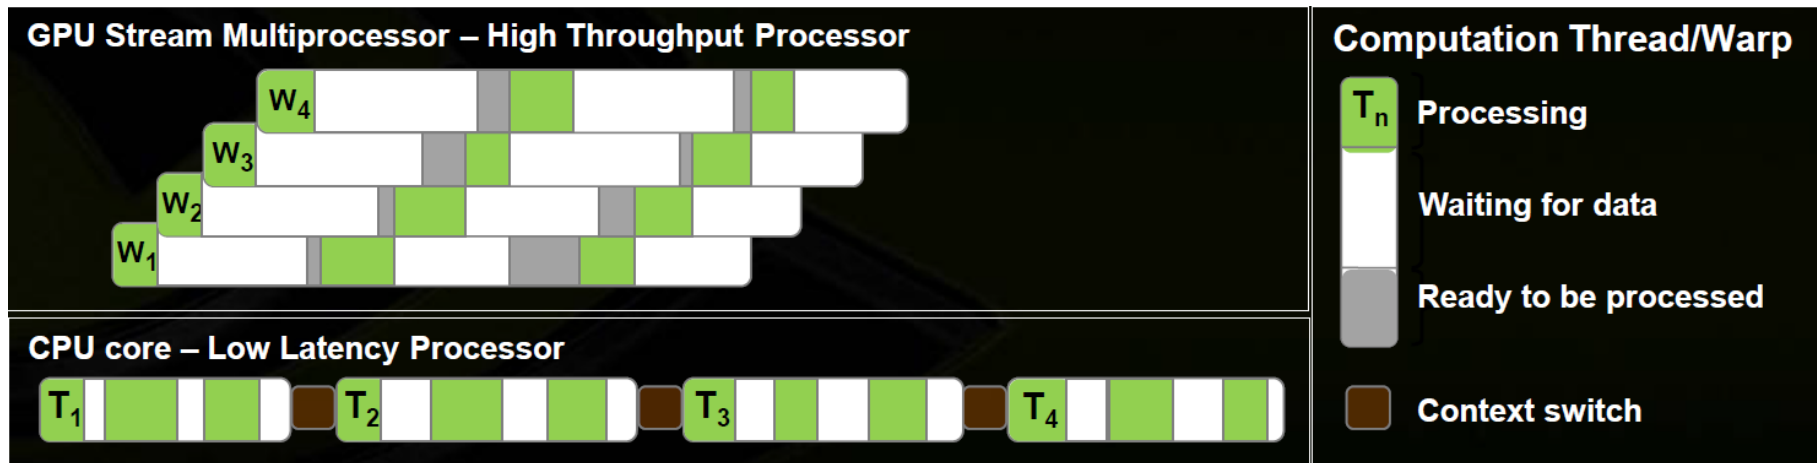
• GPU

- Optimized for data-parallel, throughput computation
- Architecture tolerant of memory latency
- More transistors dedicated to computation



CPU-GPU Cooperation (2/2)

- CPU architecture must minimize latency within each thread
- GPU architecture hides latency with computation from other thread warps



<http://www.cc.gatech.edu/~vetter/keeneland/tutorial-2011-04-14/02-cuda-overview.pdf>

OpenCL Specification

- 1. OpenCL application**
- 2. Platform model**
 - Define the host and devices.
- 3. Execution model**
 - Defines the OpenCL environment on the host and how kernels are executed on the device.
- 4. Memory model**
 - Defines the abstract memory hierarchy that kernels use
- 5. Programming model**
 - Defines how the concurrency model is mapped to physical hardware.

OpenCL Application (1/2)

OpenCL Code Consists of Two Part:

Host code:

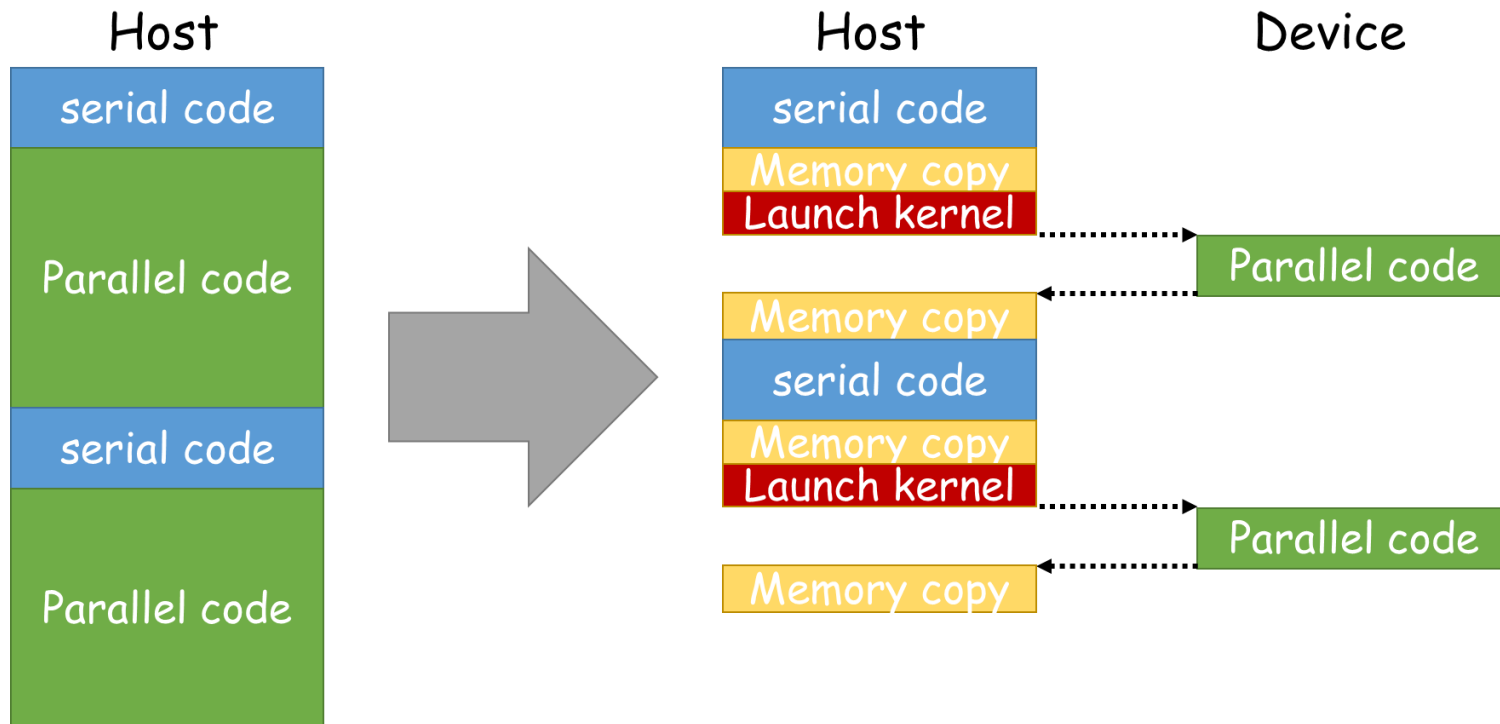
- Written in regular C or C++
- Creating the data structures that manage the host-device communication
- Run on the host (like CPU)

Kernel code:

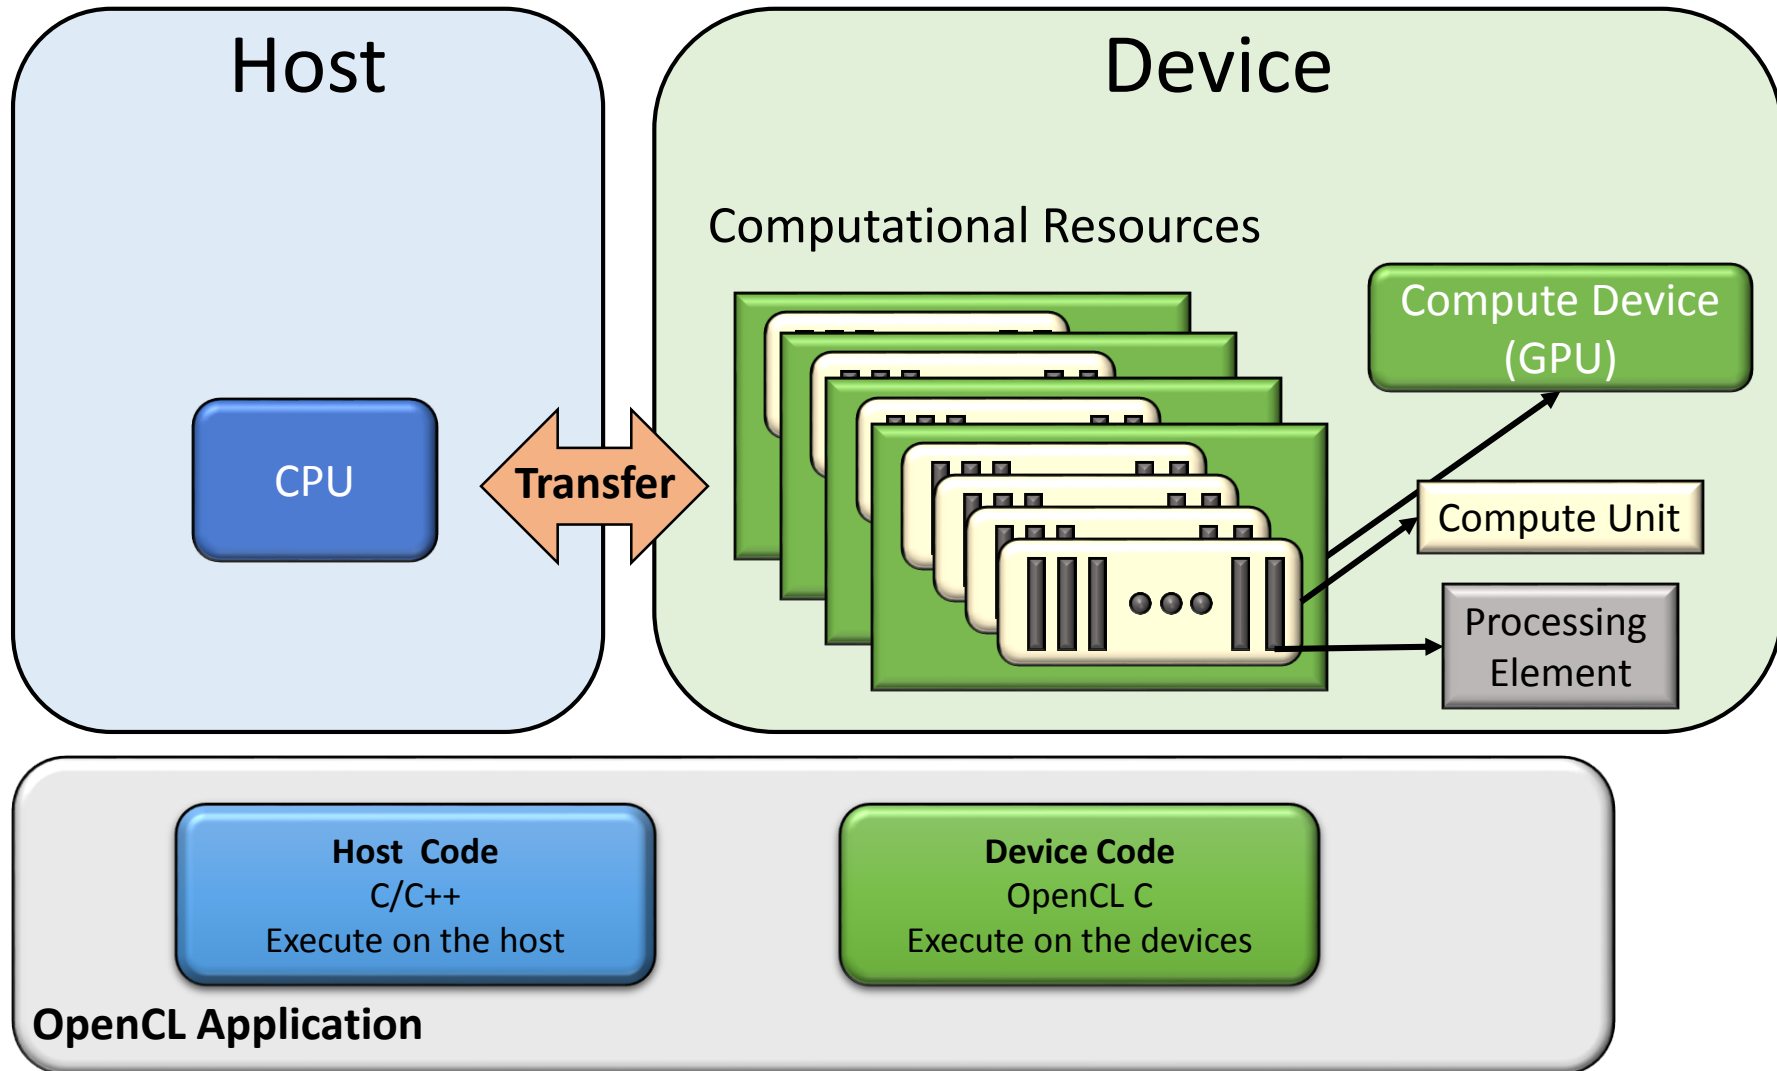
- Using the high-performance capabilities defined in the OpenCL standard.
- Run on the devices (like GPU)

OpenCL Application (2/2)

- **Serial code** executes in a **Host (CPU)** thread.
- **Parallel code** executes in many **Device (GPU)** threads across multiple processing elements.



OpenCL Platform Model



OpenCL Execution Model (1/3)

Decompose Task Into *work-items*

- Define N-dimensional computation domain
- Execute a *kernel* at each point in computation domain

• Traditional loop as function in C

```
void trad_mul( int n,  
               const float *a,  
               const float *b,  
               float *c)  
{  
    int i;  
    for (i=0; i<n; i++)  
        c[i] = a[i] * b[i];  
}
```



• OpenCL C code

```
__kernel void  
dp_mul( __global const float *a,  
         __global const float *b,  
         __global float *c)  
{  
    int id = get_global_id(0);  
    c[id] = a[id] * b[id];  
} // execute over n "work items"
```


OpenCL Execution Model (2/3)

An N-dimension Domain of Work-items

- Kernels are executed across a **global domain** of *work-items*
- Work-items are grouped into **local domain** of *work-groups*

Global dimensions: 1024 x 1024

Local dimensions: 128 x 128

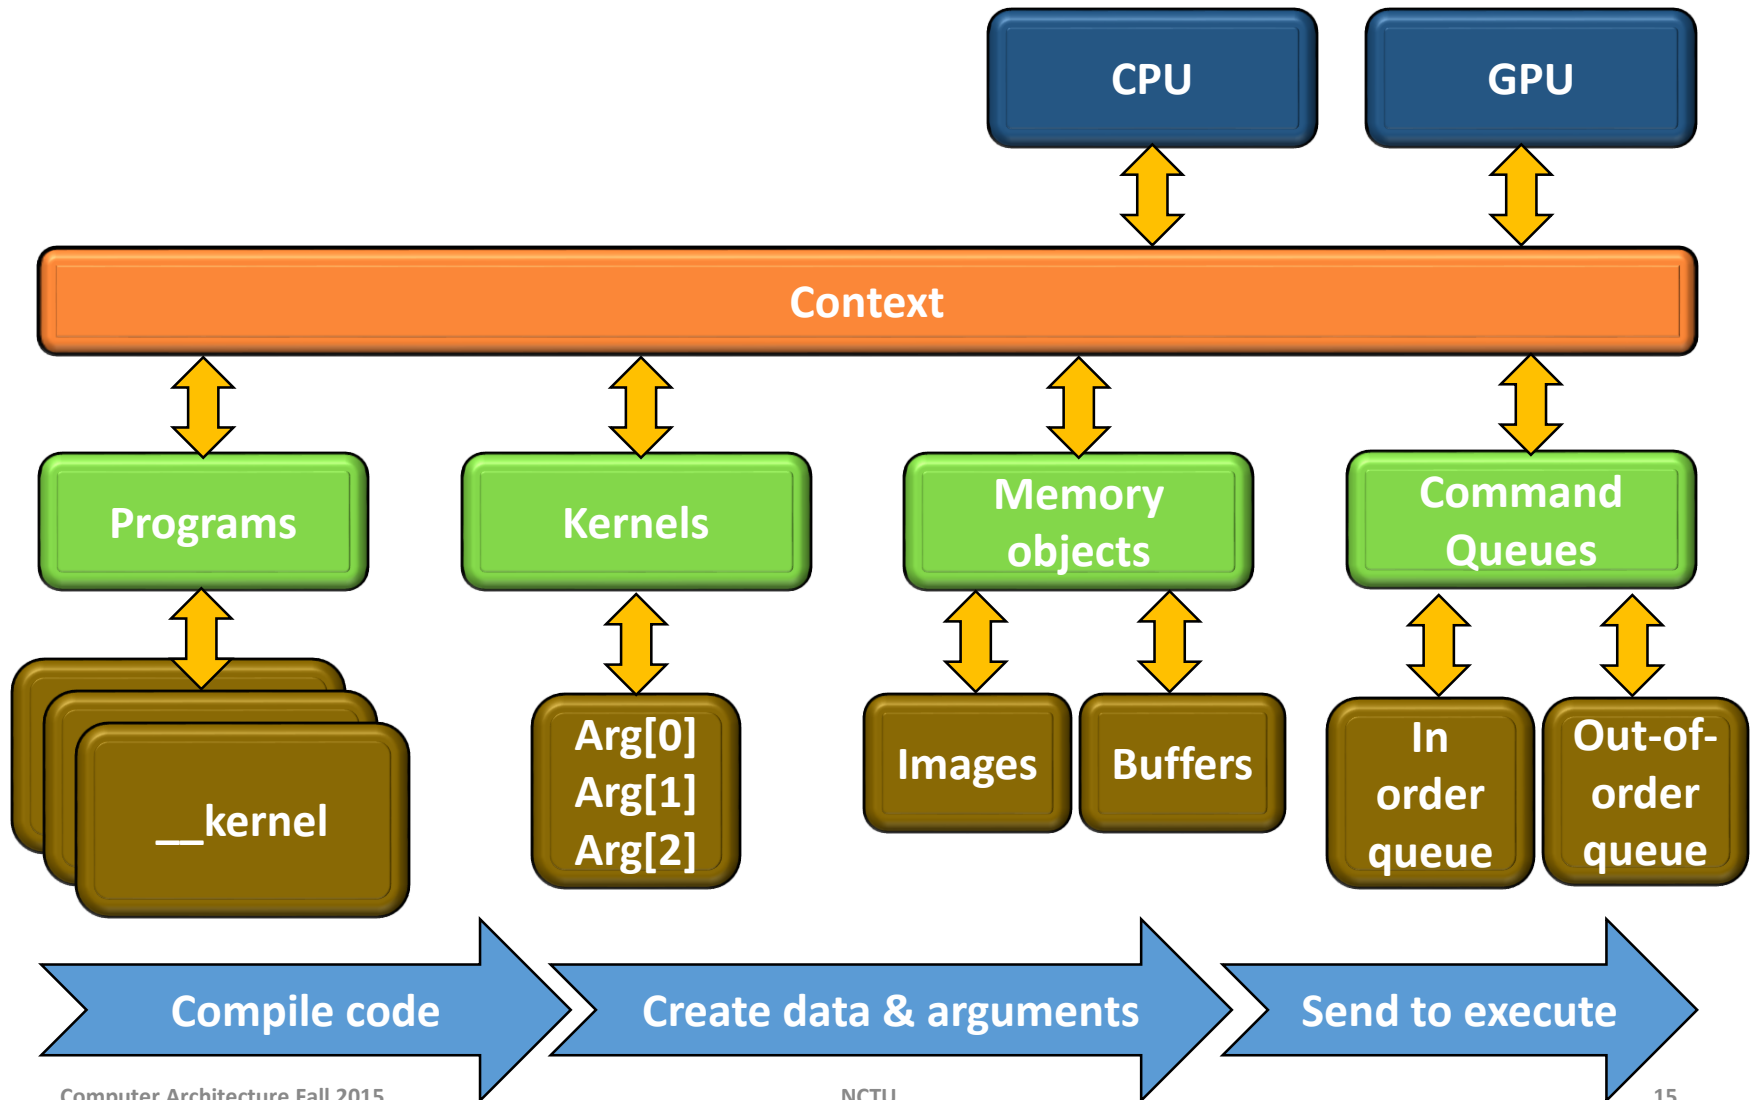


OpenCL Execution Model (3/3)

The application runs on **the host** which submits task to **the devices**

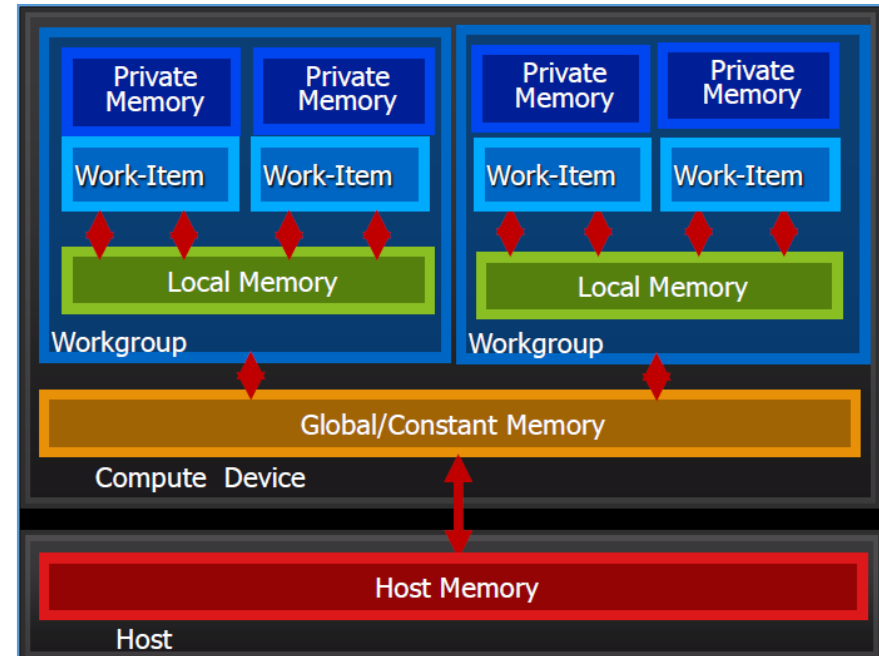
- **Work-item:**
the basic unit of work on the OpenCL device
- **Kernel:**
the code for a work-item (basically a C function)
- **Program:**
Collection of kernels and other functions

OpenCL Framework



OpenCL Memory Model

- **Private Memory**
 - Per work-item
- **Local Memory**
 - Shared within a workgroup
- **Global/Constant Memory**
 - Visible to all workgroups
- **Host Memory**
 - On the CPU



OpenCL Programming Flow (1/5)

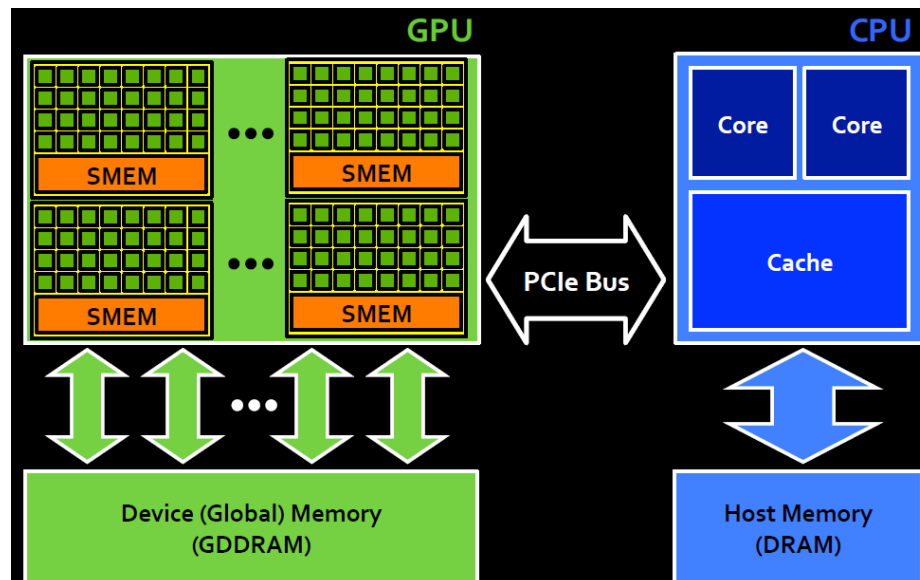
Standard Execute Flow

Step 1: Copy data to GPU memory

Step 2: Launch the kernels on GPU

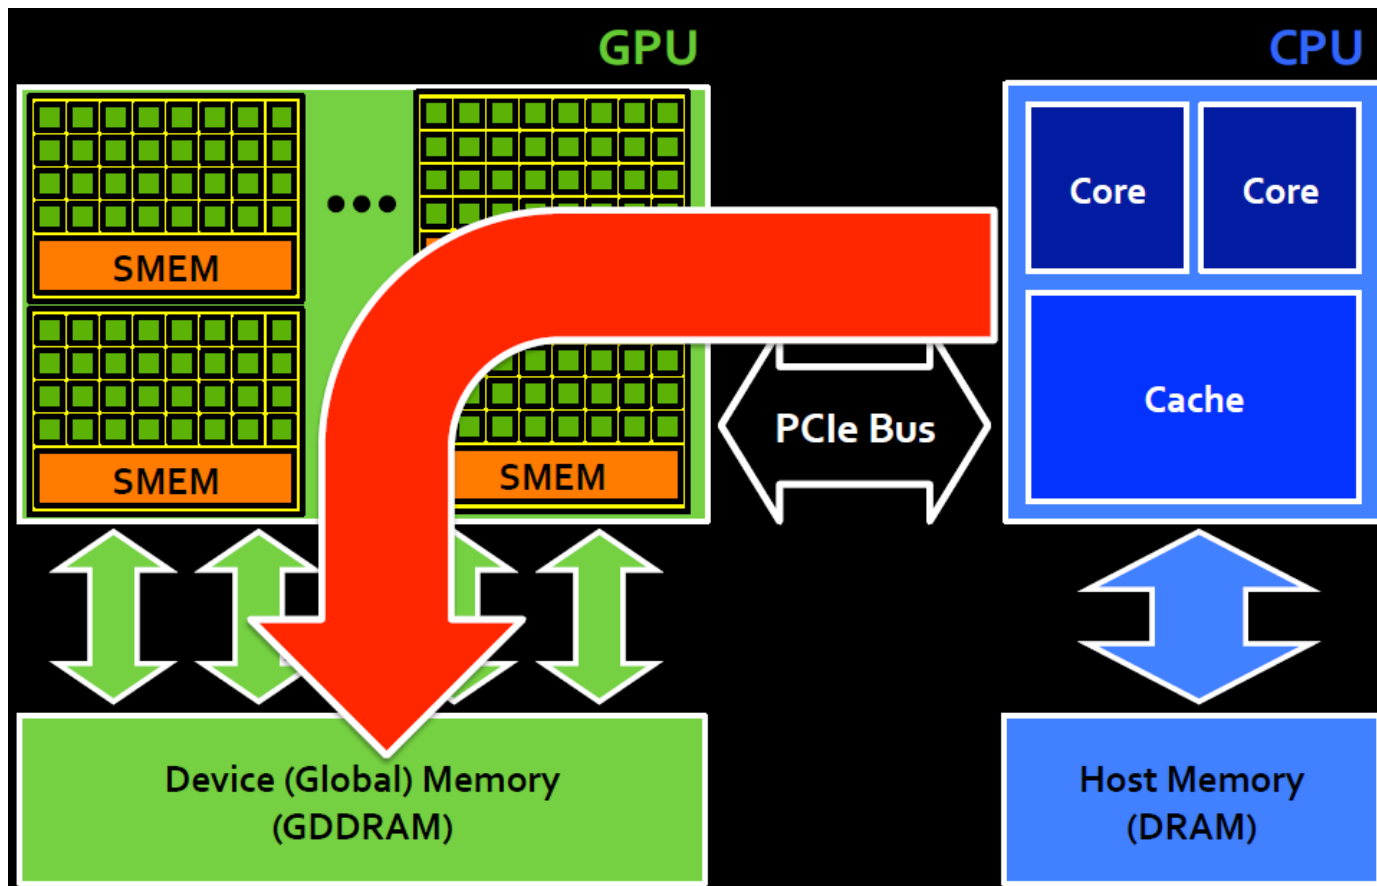
Step 3: Execute kernels on GPU

Step 4: Copy data to CPU memory



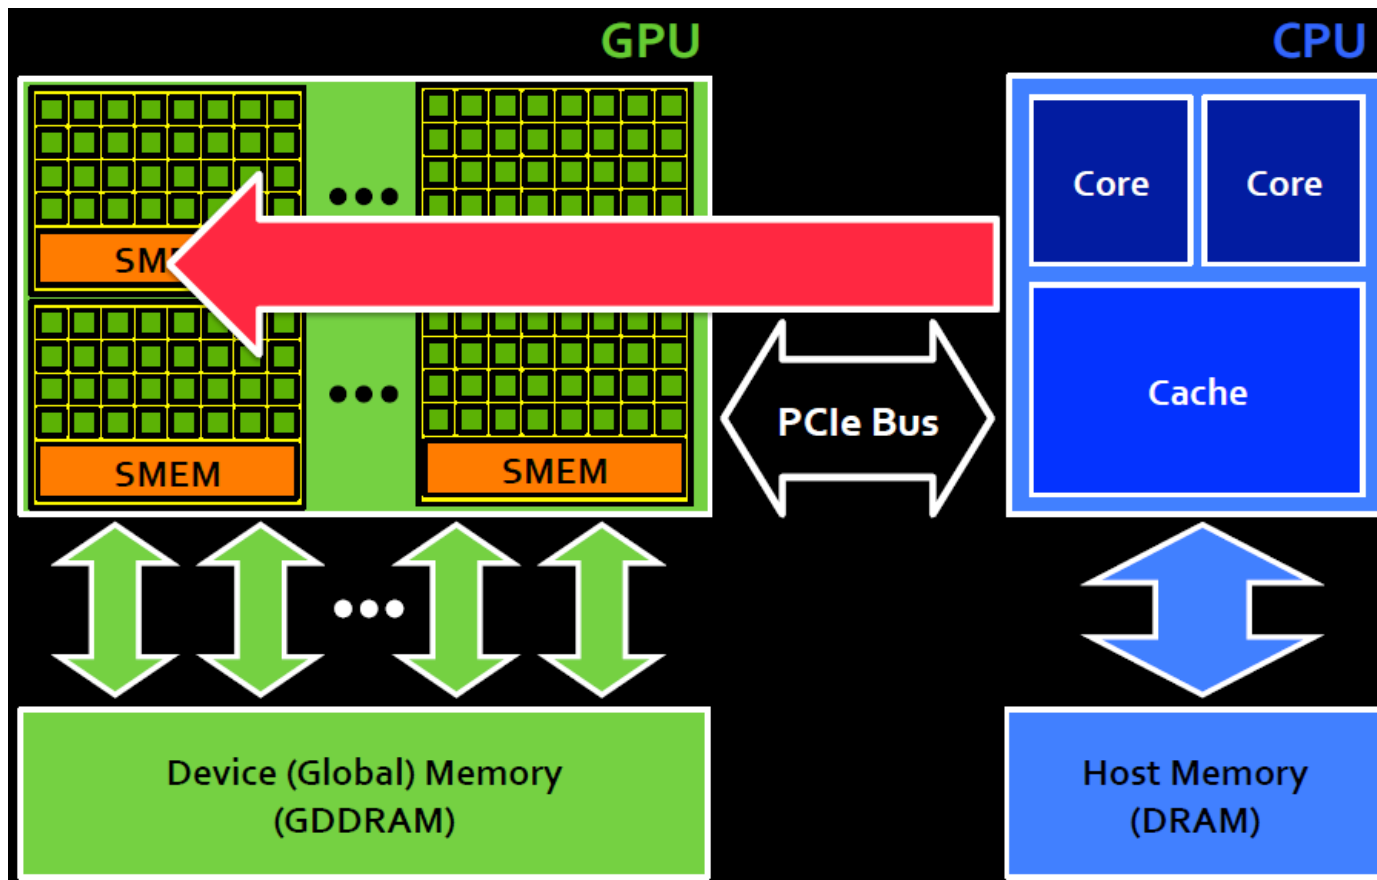
OpenCL Programming Flow (2/5)

Step 1: Copy Data to GPU Memory



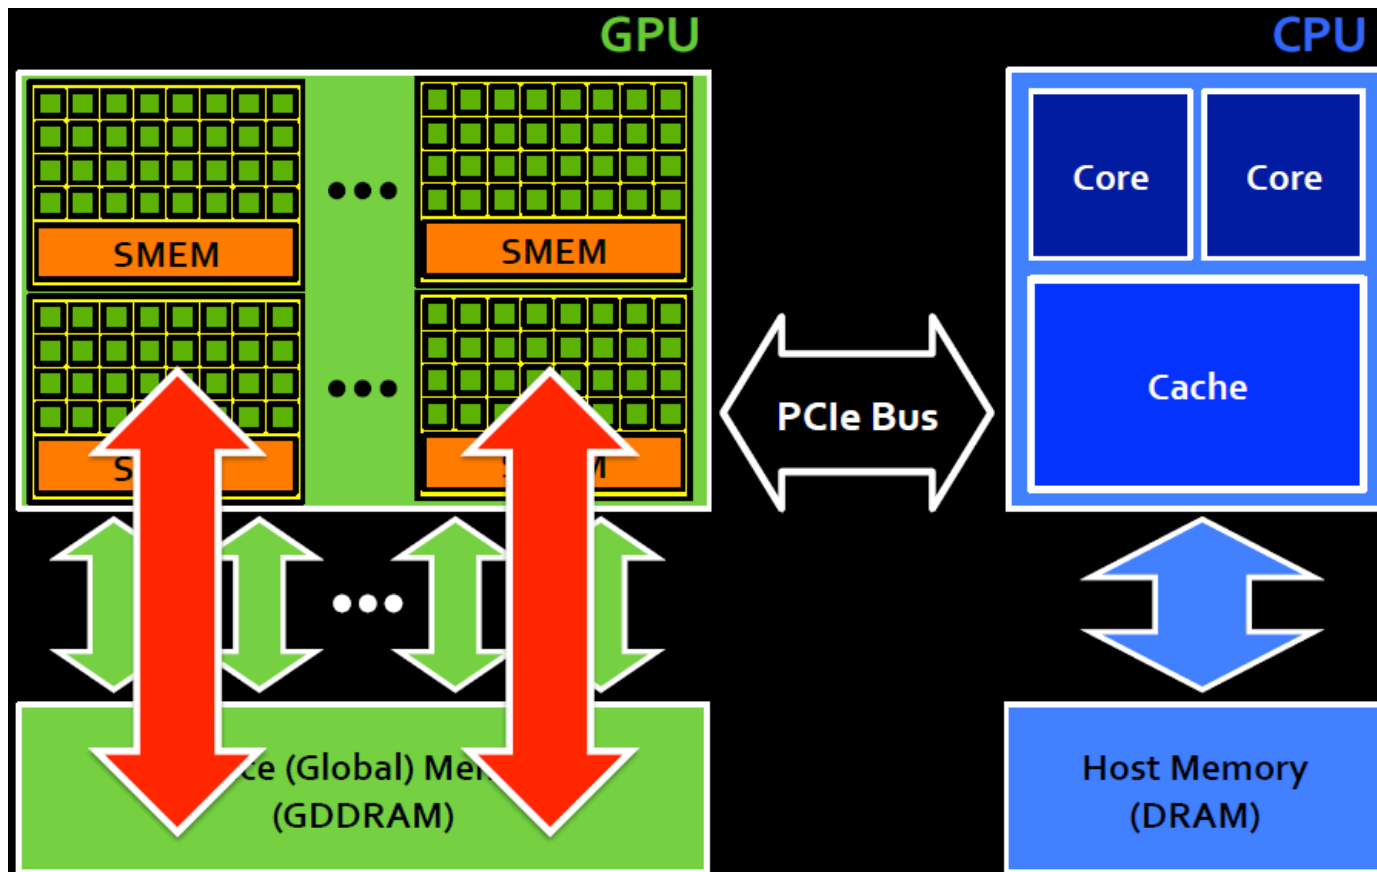
OpenCL Programming Flow (3/5)

Step 2: Launch The Kernels on GPU




OpenCL Programming Flow (4/5)

Step 3: Execute Kernels on GPU



OpenCL Programming Flow (5/5)

Step 4: Copy Data to CPU Memory

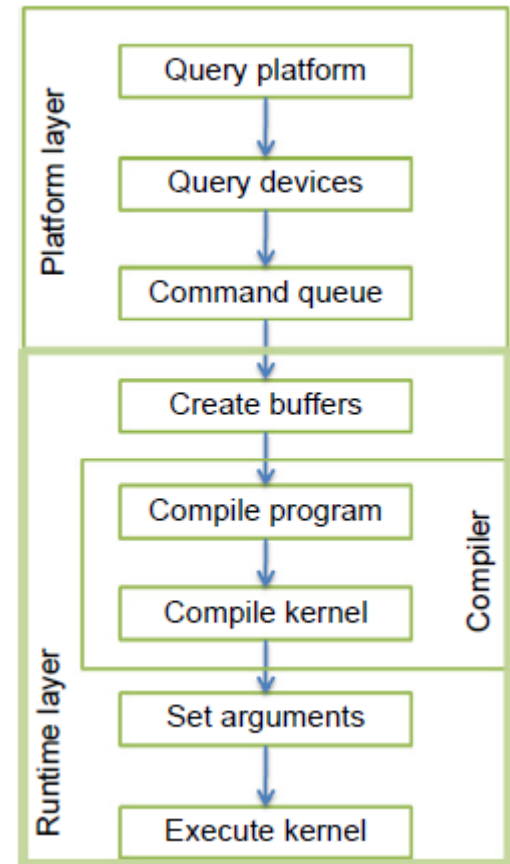


Outline

- **Introduction to OpenCL**
- **OpenCL Framework**
 - Platform layer
 - Compiler for OpenCL C
 - Runtime
- **The Flow of Host Program**
 - Example: Hello world

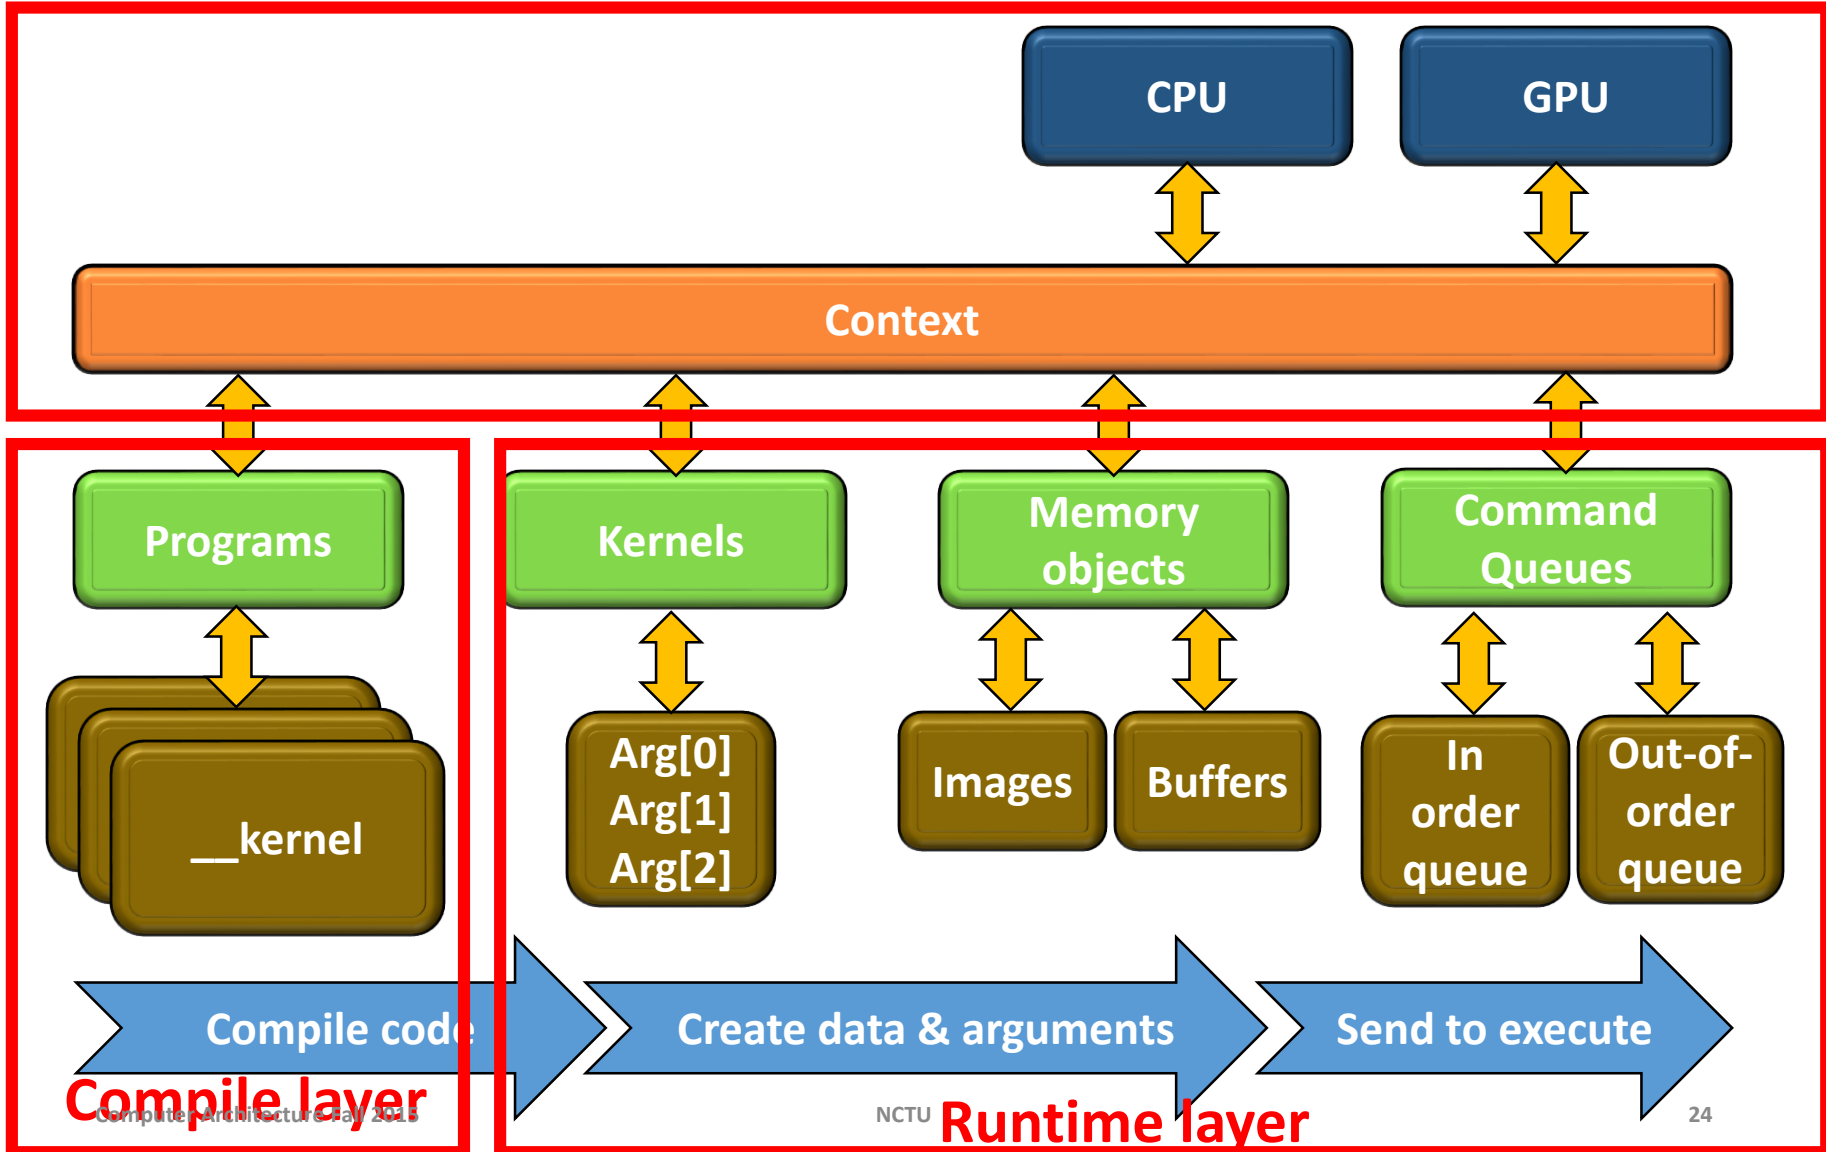
OpenCL Framework

- **Platform layer**
 - Platform query and context creation
- **Compiler for OpenCL C**
- **Runtime**
 - Memory management and command execution within a context



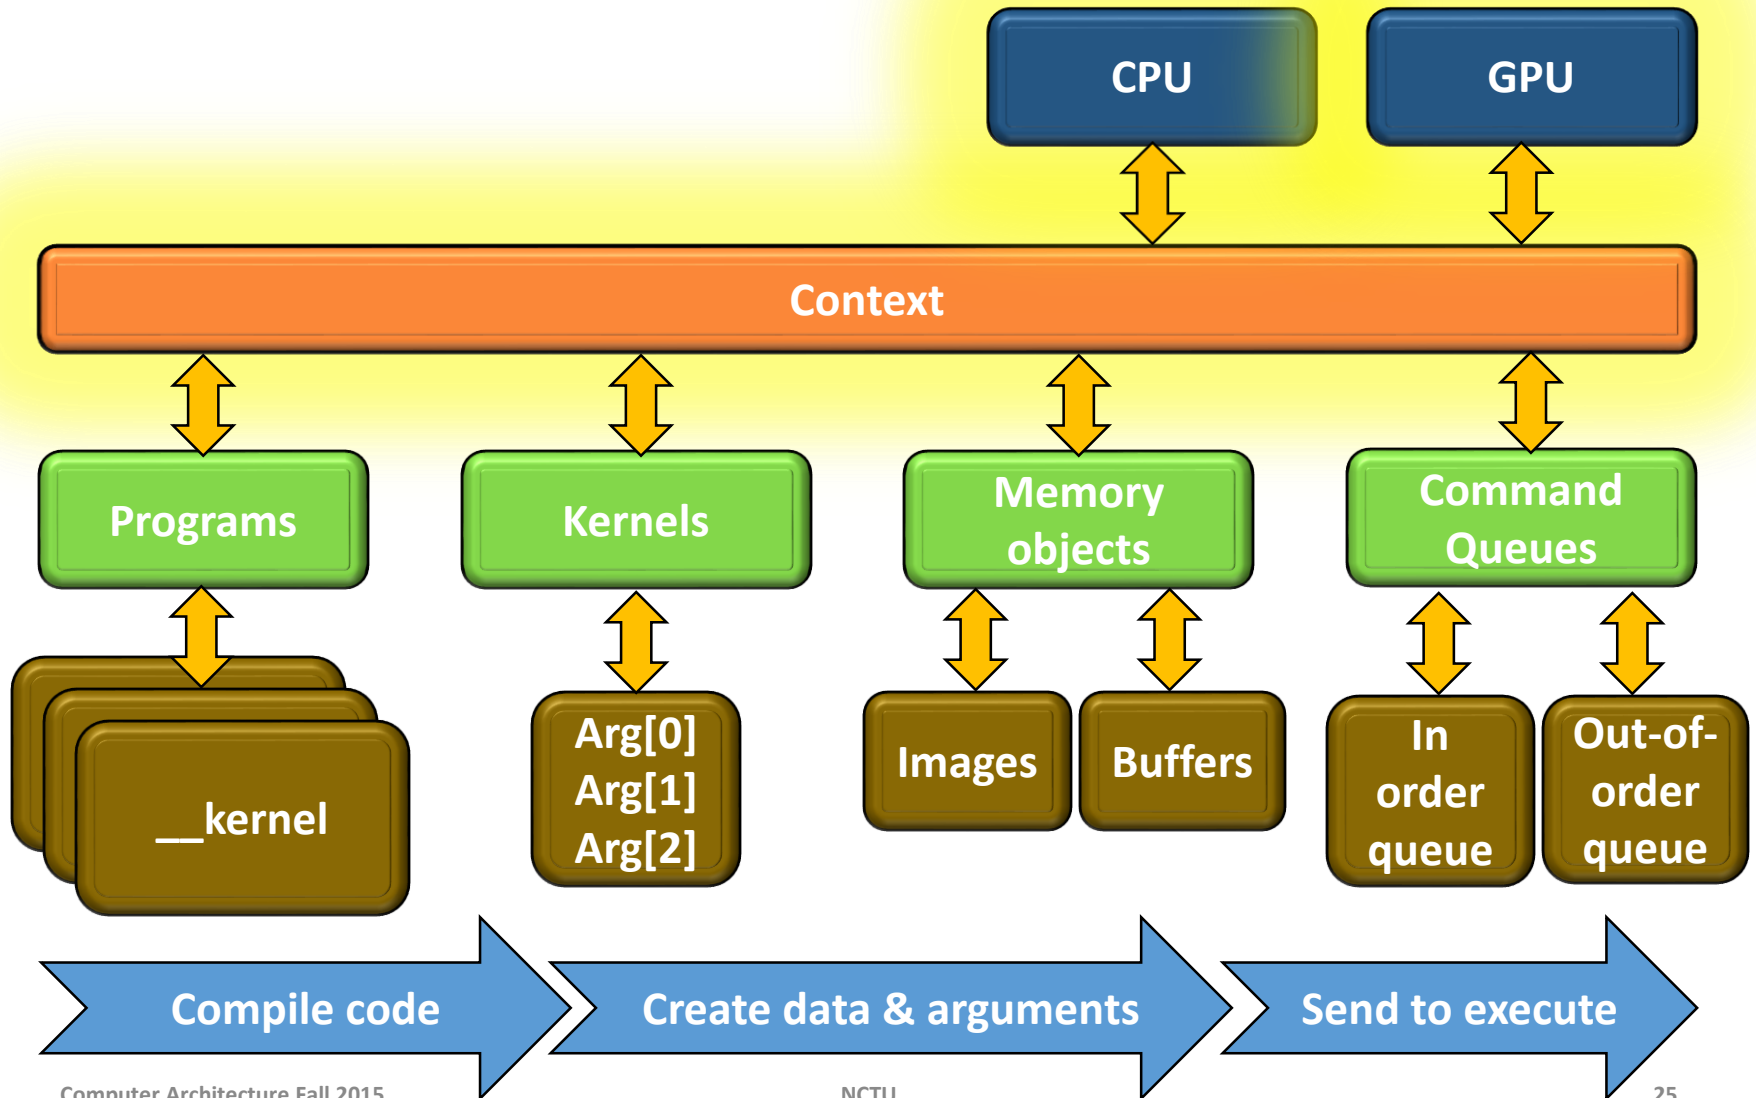
OpenCL Framework

Platform layer




OpenCL Framework

Platform layer



OpenCL Framework

Platform layer




Create Platform

- ***clGetPlatformIDs()***: Obtain the list of platforms available.
- ***clGetPlatformInfo()***: Name, version, vendor, extensions
- ***clGetDeviceIDs()***:
Obtain the list of devices available on a platform.
- ***clGetDeviceInfo()***: Name, type, capabilities

OpenCL Framework

Platform layer



Create Context

clCreateContext(): Creates an OpenCL context.

Context =
`cl_context`

{

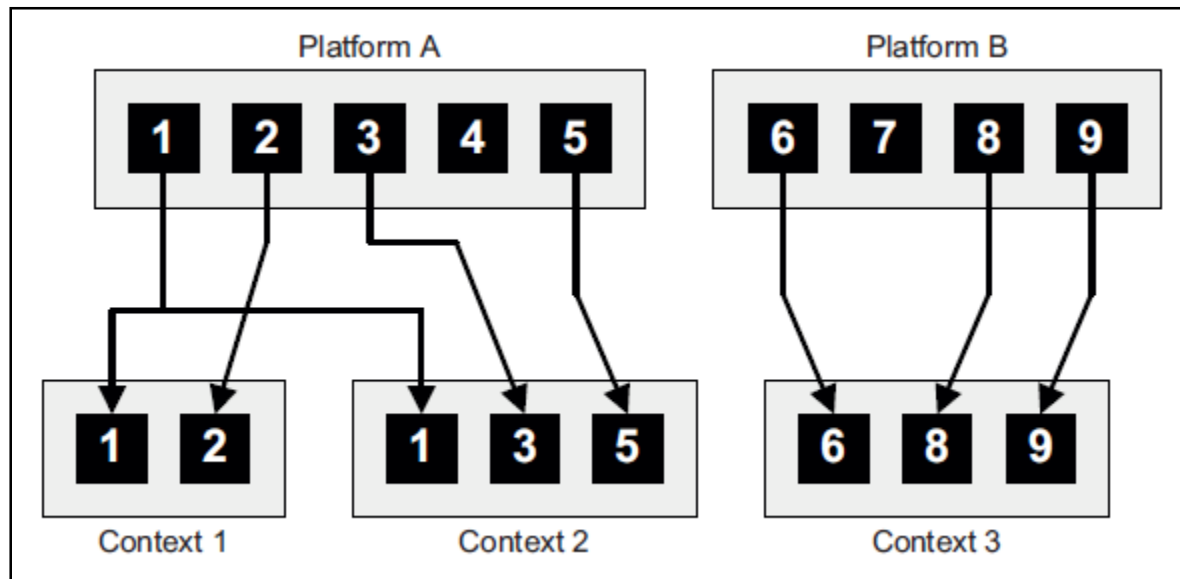
- One or more devices
`cl_device_id`
- Memory and device code
`cl_mem` `cl_program`
- Command queue to send command to devices
`cl_command_queue`

OpenCL Framework

Platform layer

Create Context

- The devices in a context must be provided by the same platform.



OpenCL Framework

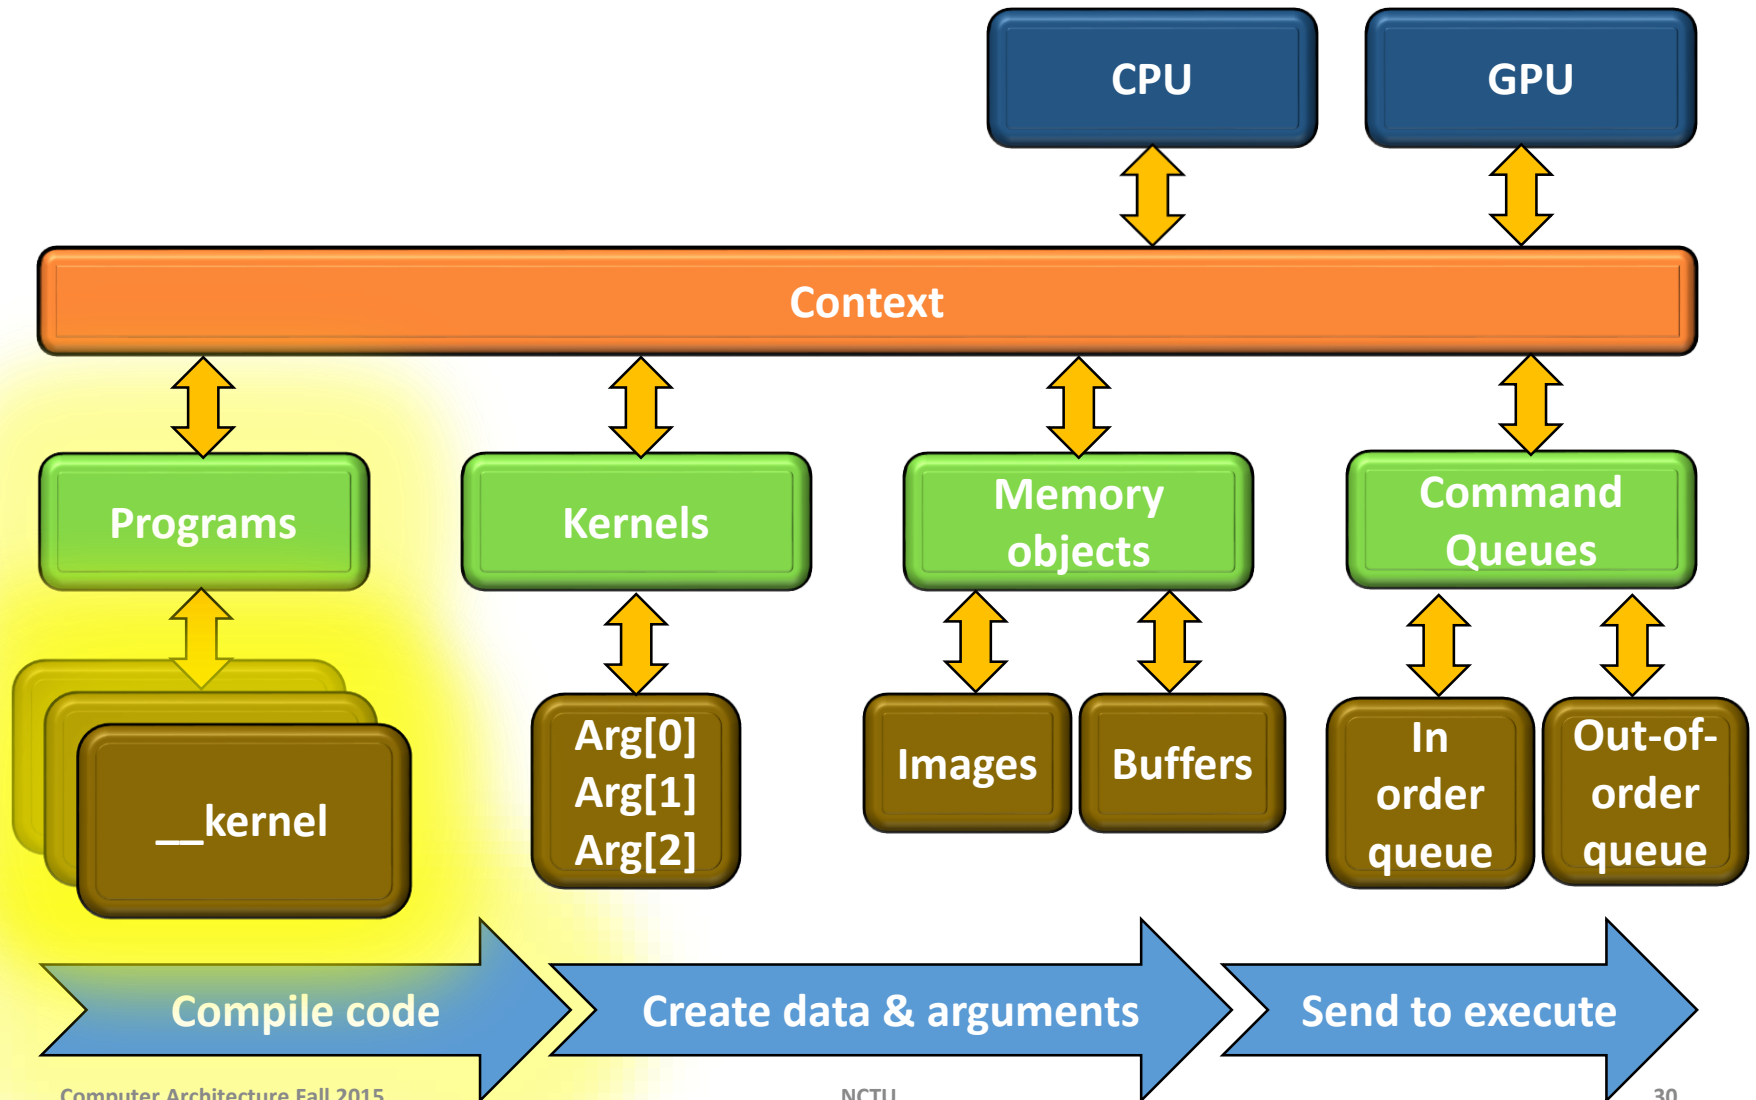
Platform layer

Simple Example: Context Creation

```
// Get the platform ID
cl_platform_id platforms;
clGetPlatformIDs(1, &platforms, NULL);
// Get the first GPU device associated with the platform.
cl_device_id device;
clGetDeviceIDs(platforms, CL_DEVICE_TYPE_ALL, 1, &device, NULL);
// Create an OpenCL context for GPU device
cl_context context;
context=clCreateContext(NULL, 1, devices, NULL, NULL, NULL);
```

OpenCL Framework

Compile OpenCL C




OpenCL Framework

Compile OpenCL C

OpenCL C

- **Derived from ISO C99 (with some restrictions)**
- **Language Features Added**
 - Work-items and work-groups
 - Vector types
 - Synchronization
 - Address space qualifiers
- **Includes some built-in functions**
 - Image manipulation
 - Work-item manipulation
 - Math functions




OpenCL Framework

Compile OpenCL C

OpenCL C Optional Extensions

- **Extensions are optional features exposed through OpenCL**

- Double precision
- Atomic functions
- Byte-addressable stores
- Print functions

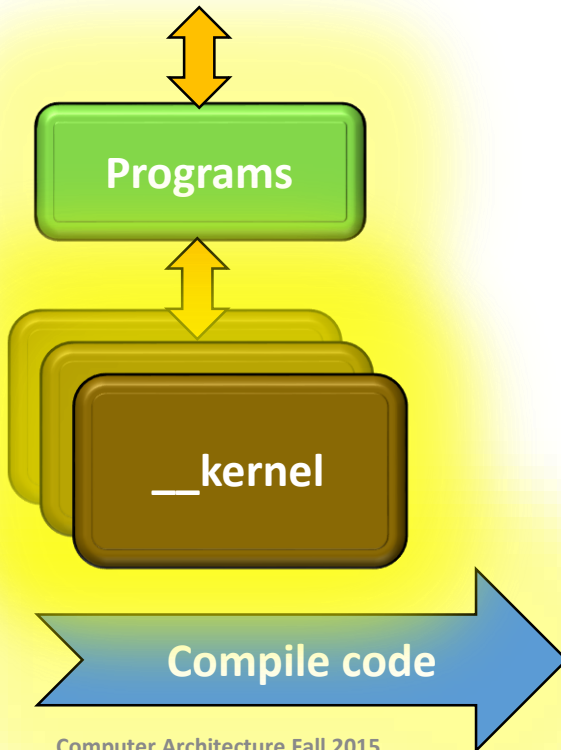


OpenCL Framework

Compile OpenCL C

Language Features

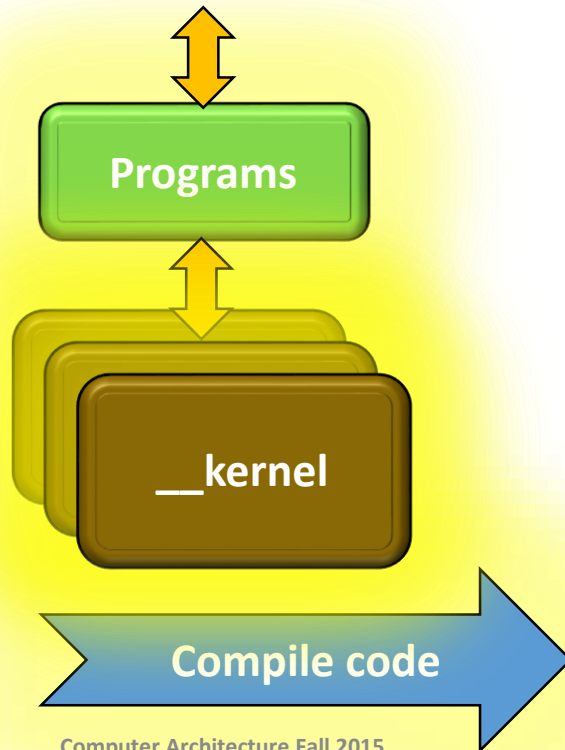
- Work-items and work-groups
- Vector types
- Address space qualifiers
- Synchronization



OpenCL Framework

Compile OpenCL C

OpenCL C Language Restrictions




- Pointers to functions are not allowed
- Pointers to pointers allowed within a kernel, but not as an argument
- Bit-fields are not supported
- Variable-length arrays and structures are not supported
- **Recursion is not supported**
- **Double types are not supported, but reserved**

OpenCL Framework

Compile OpenCL C

OpenCL Program Structure


- OpenCL C code is called a program (*cl_program*).
- A program is a collection of functions called kernels (*cl_kernel*).
- Programs are compiled at runtime through a series of API calls.
- This runtime compilation gives the system an opportunity to optimize for a specific device.



OpenCL Framework

Compile OpenCL C

Device Code Compilation And Execution




- A ***cl_program*** object encapsulates some source code (kernel functions) and its last successful build.
 - *clCreateProgramWithSource()* // Create program from source
 - *clBuildProgram()* // Compile program
- A ***cl_kernel*** object encapsulates the values of the kernel's arguments used when the kernel is executed
 - *clCreateKernel()* // Create kernel from successfully compiled program
 - *clSetKernelArg()* // Set values of kernel's arguments

OpenCL Framework

Compile OpenCL C

The Process of Creating a Kernel

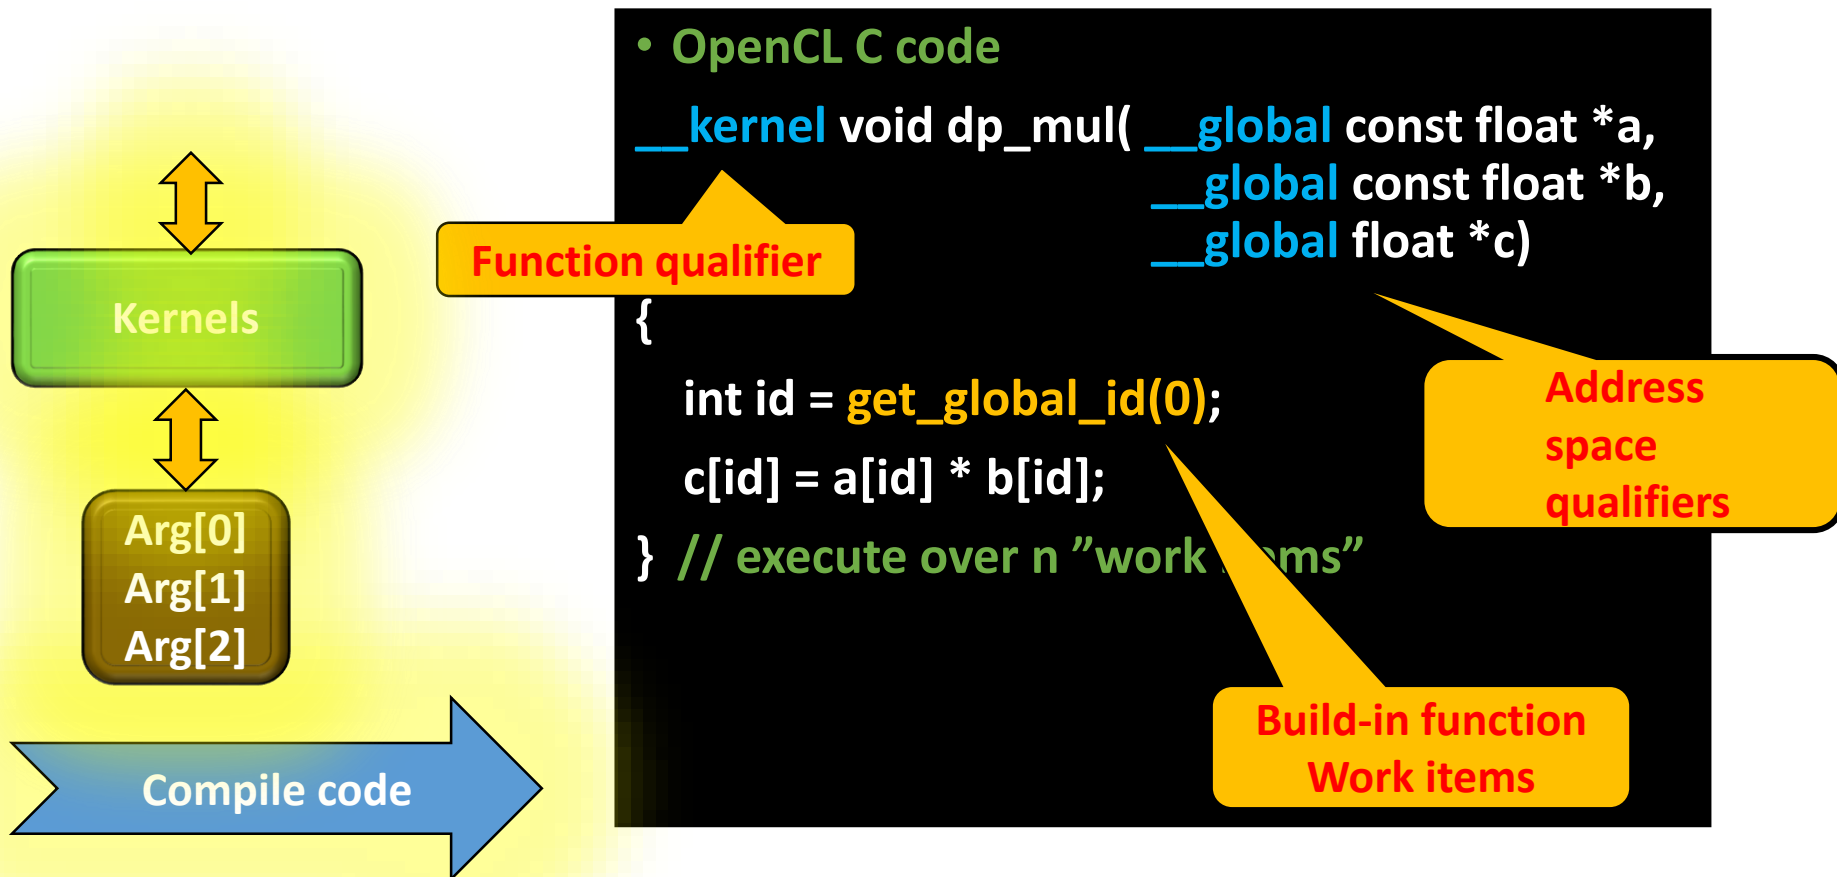


1. The OpenCL C source code is stored in a character string.
2. The source code is turned into a program object, *cl_program*, by calling *clCreateProgramWithSource()*.
3. The program object is then compiled with *clBuildProgram()*. If there are compile errors, they will be reported here.

OpenCL Framework

Compile OpenCL C

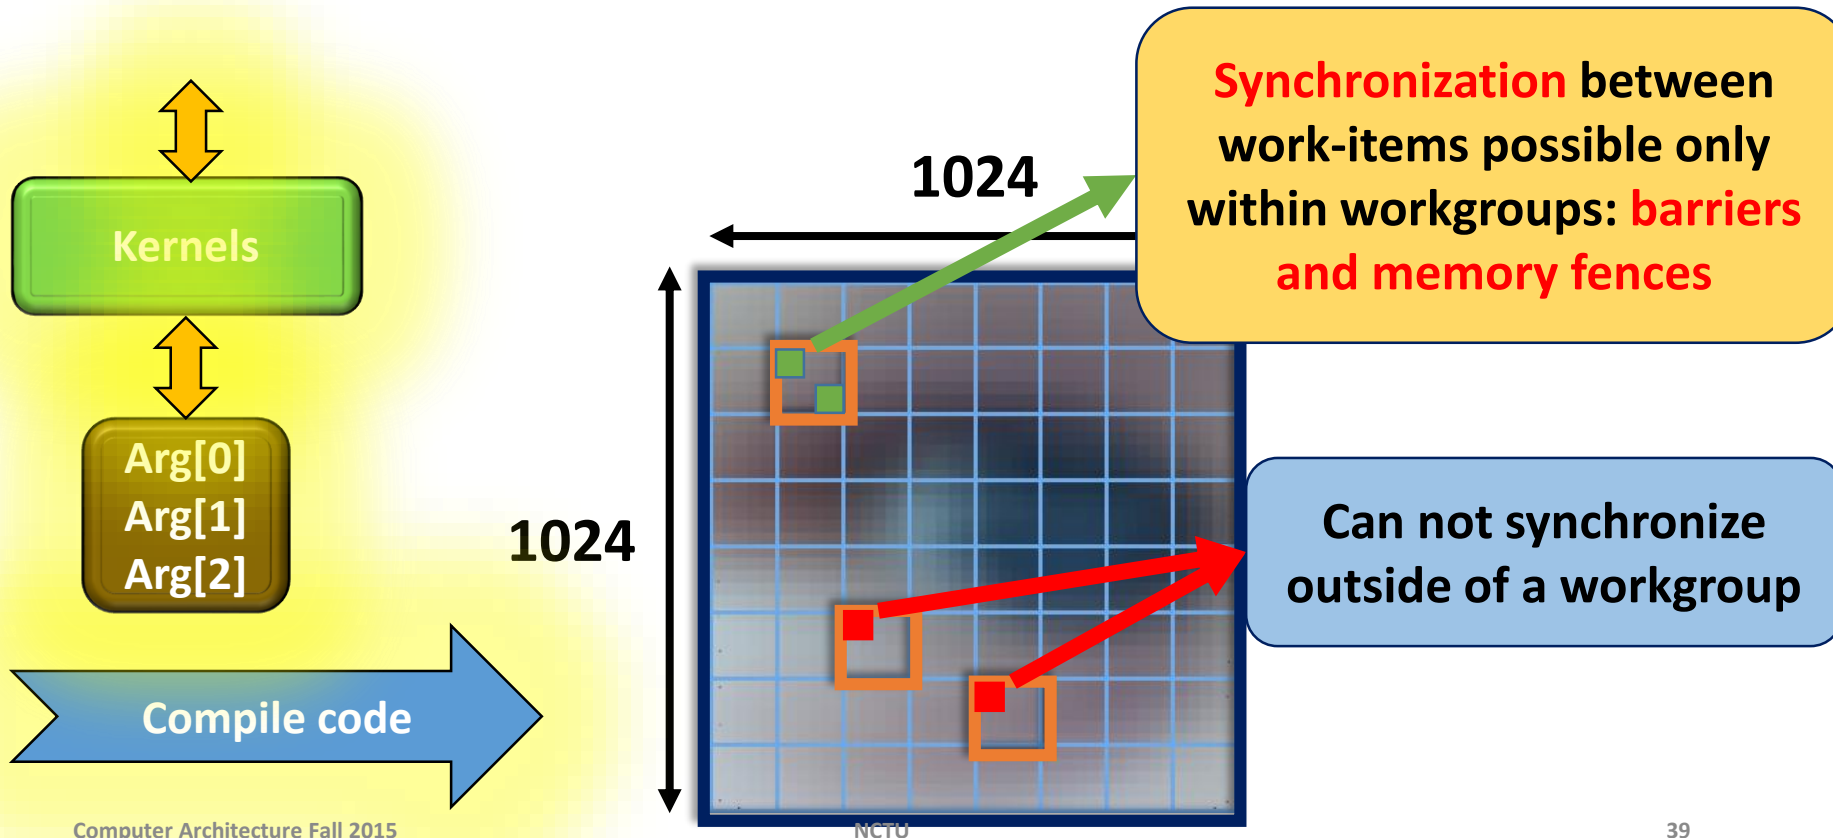
Simple Example: Language Features



OpenCL Framework

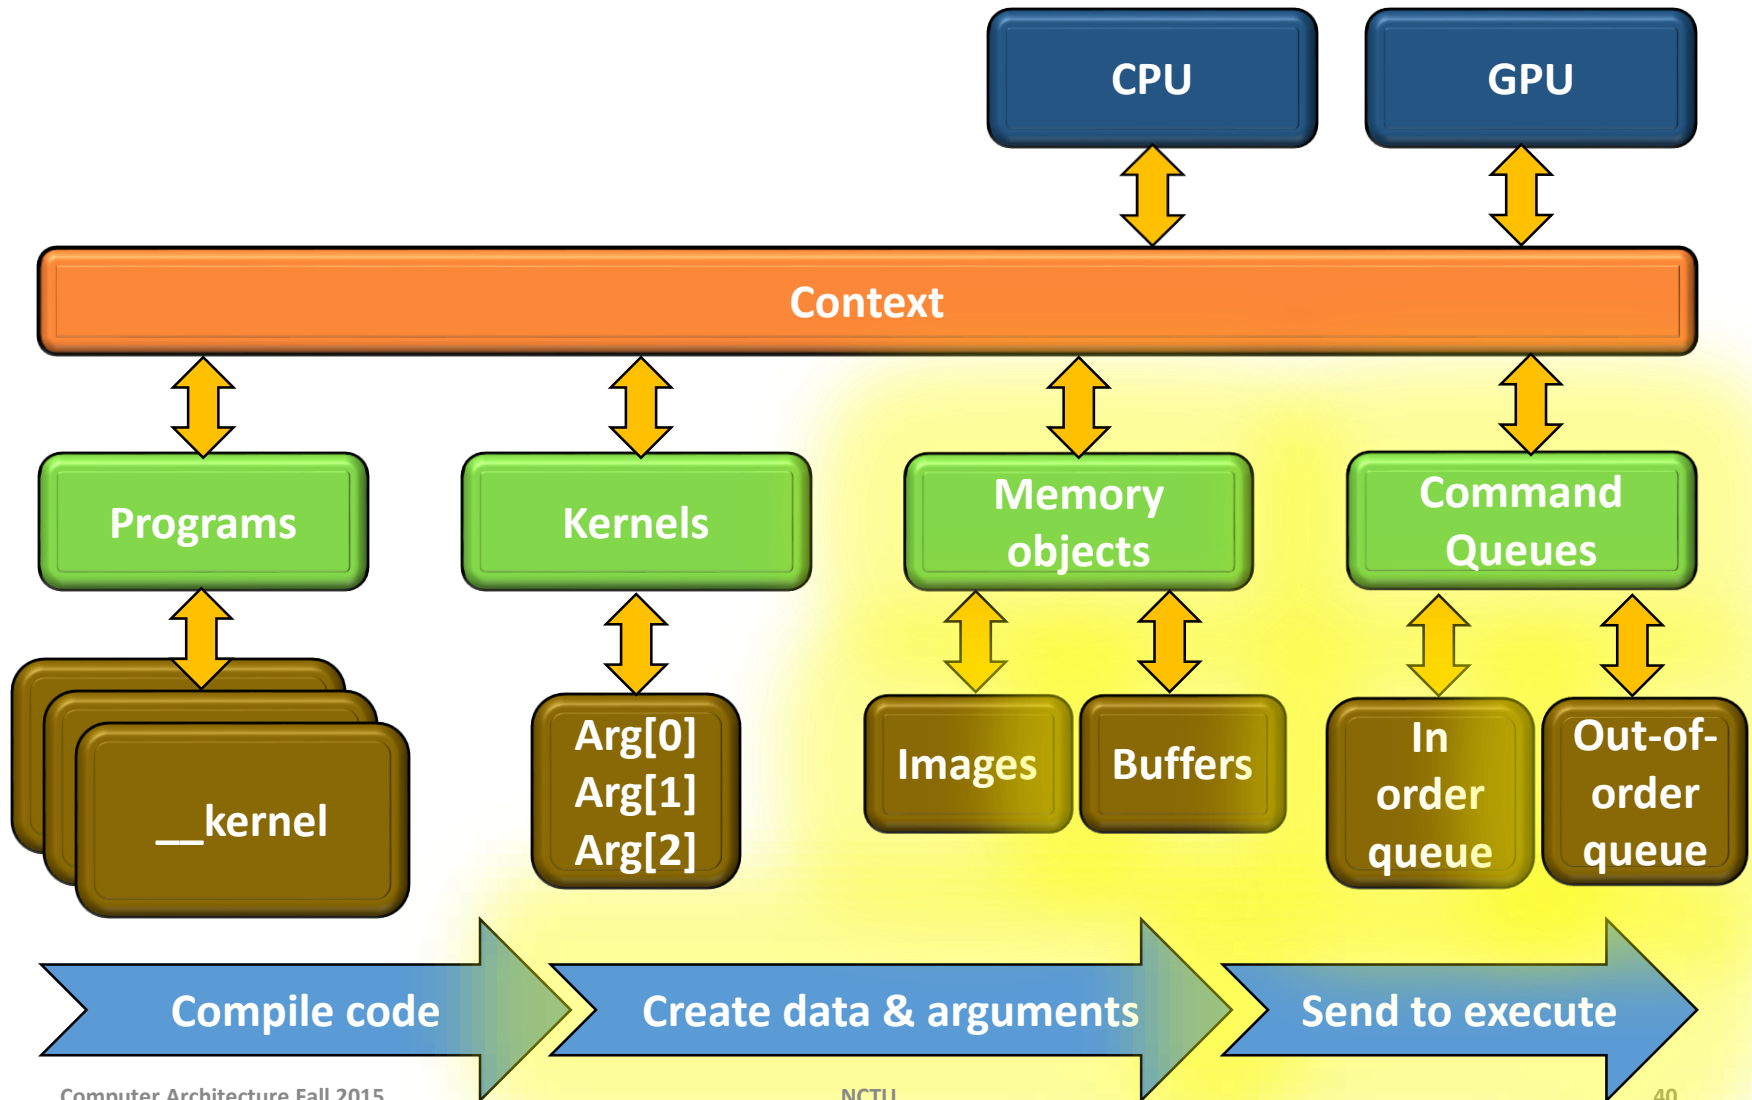
Compile OpenCL C

Simple Example: Language Features



OpenCL Framework

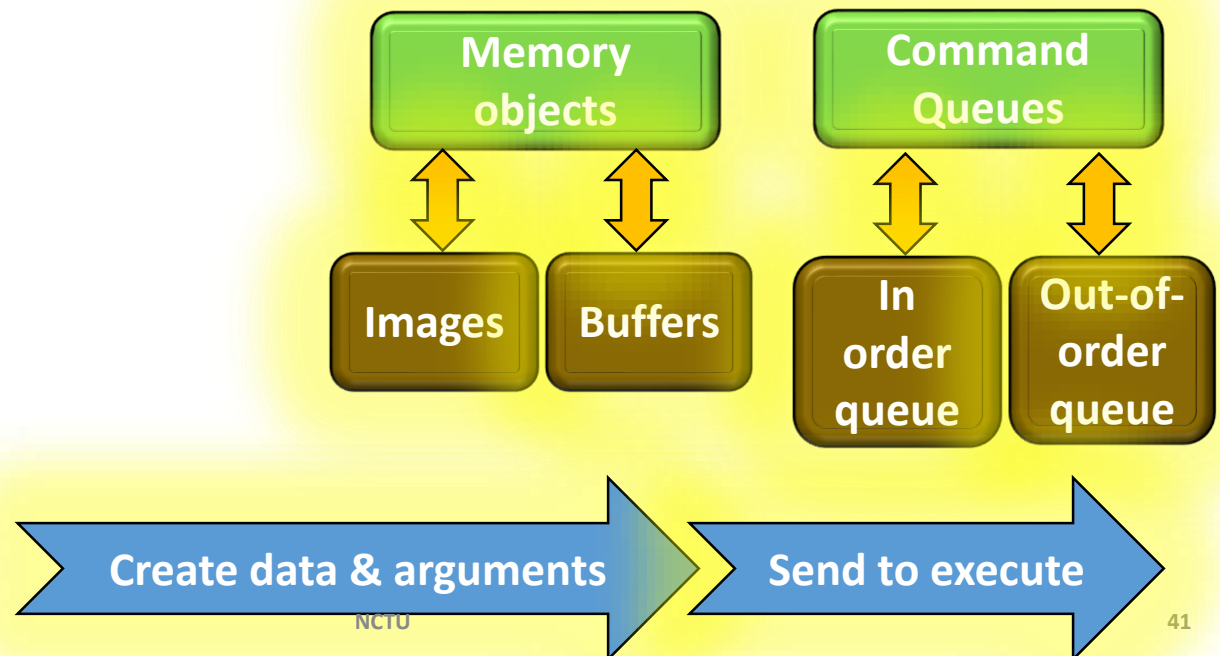
Runtime layer



OpenCL Framework

Runtime layer

- Device memory allocation and management
- Command queues creation and management
- Device code compilation and execution
- Event creation and management (synchronization, profiling)




OpenCL Framework

Runtime layer

Device Memory Allocation And Management

- When the device performs a task, you have to provide at least three pieces of information:
 1. The instructions to be executed,
 2. A buffer containing data to be processed
 3. A buffer where processed data should be stored
- In OpenCL, memory objects serve as standard packages for transfer data between a host and device.



OpenCL Framework

Runtime layer

Device Memory Allocation And Management

- OpenCL defines two types of memory objects(**cl_mem**):

- **Buffer objects:**


One-dimensional array

*clCreateBuffer(cl_context context,
cl_mem_flags options,
size_t size,
void *host_ptr,
cl_int *error)*

CL_MEM_READ_WRITE
CL_MEM_WRITE_ONLY
CL_MEM_READ_ONLY
CL_MEM_USE_HOST_PTR
CL_MEM_COPY_HOST_PTR
CL_MEM_ALLOC_HOST_PTR

- **Image objects:**

*clCreateImage2D()
clCreateImage3D()*




OpenCL Framework

Runtime layer

Command Queue Creation

- Command queue is mechanism of communication between host and devices.
- One command queue needs to be created per device.
- Two modes of execution:
 - In-order
 - Out-of-order

• *clCreateCommandQueue()*



OpenCL Framework

Runtime layer

How to Send Task to Devices

- OpenCL provides many functions that start with “*clEnqueue**”, and each of them dispatched a command to a device through a command queue.

- Function:


clEnqueueTask

clEnqueueReadBuffer

clEnqueueWriteBuffer

clEnqueueNDRangeKernel

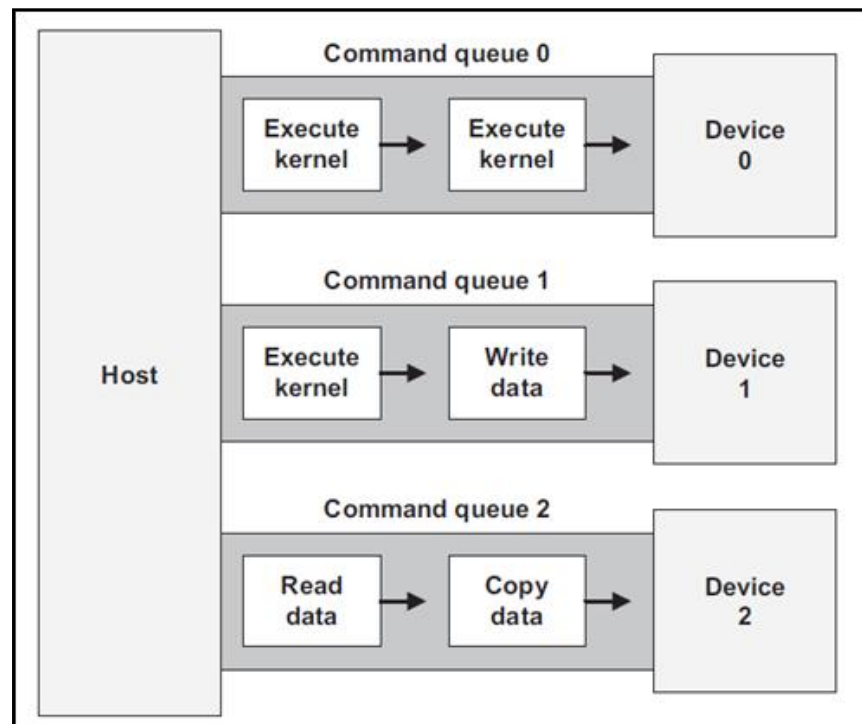
...



OpenCL Framework

Runtime layer

Command Queue Creation



OpenCL Framework

Runtime layer

How to Send Task to Devices

```
cl_int clEnqueueNDRangeKernel (cl_command_queue command_queue,  
                                cl_kernel kernel,  
                                cl_uint work_dim,  
                                const size_t *global_work_offset,  
                                const size_t *global_work_size,  
                                const size_t *local_work_size,  
                                cl_uint num_events_in_wait_list,  
                                const cl_event *event_wait_list,  
                                cl_event *event)
```

NDRange can be 1-,
2-, 3-dimensions

➤ **Global_work_size**

Define the number of global work-items in work_dim dimensions

➤ **Local_work_size**


Define the number of work-items in work_dim dimensions make up a work-group

OpenCL Framework

Runtime layer

Example: 2-dimensions

- Global dimensions: 1024 x 1024
- Local dimensions: 128 x 128



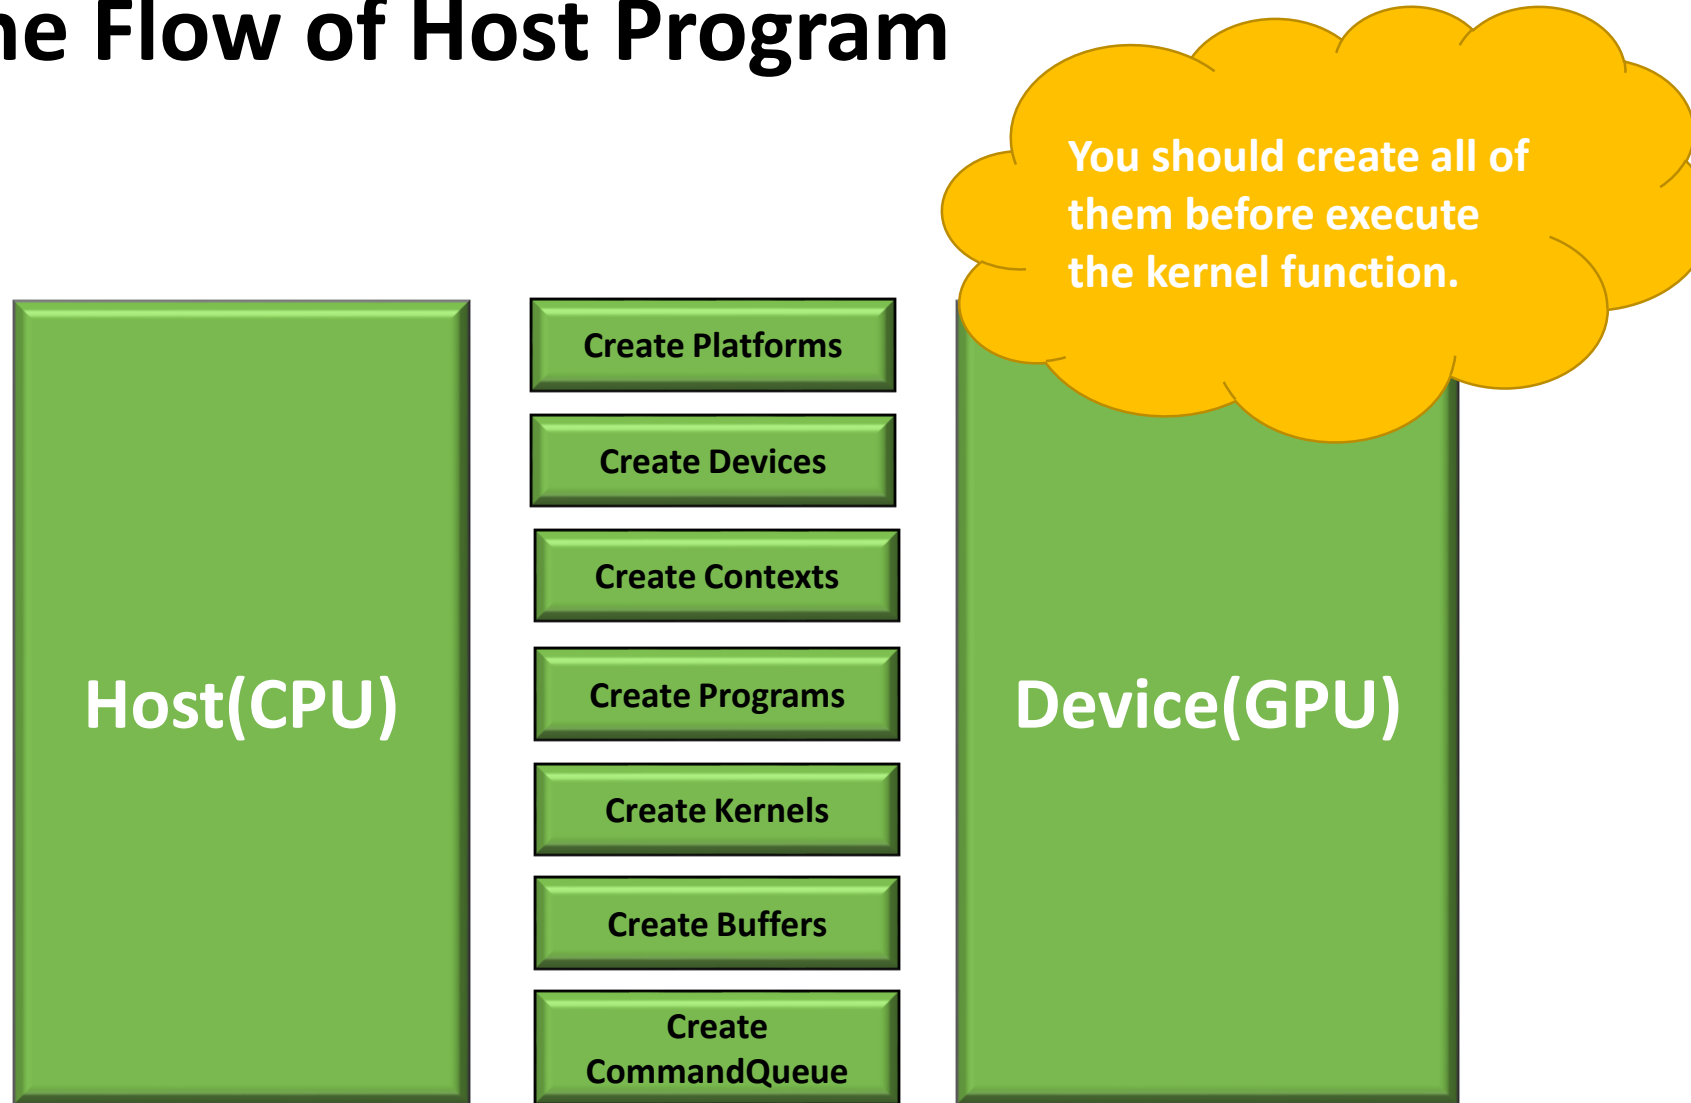
Synchronization between work-items possible only within workgroups: **barriers** and **memory fences**

Can not synchronize outside of a workgroup

Outline

- **Introduction to OpenCL**
- **OpenCL Framework**
 - Platform layer
 - Compiler for OpenCL C
 - Runtime
- **The Flow of Host Program**
 - Example: Hello world

The Flow of Host Program



First program

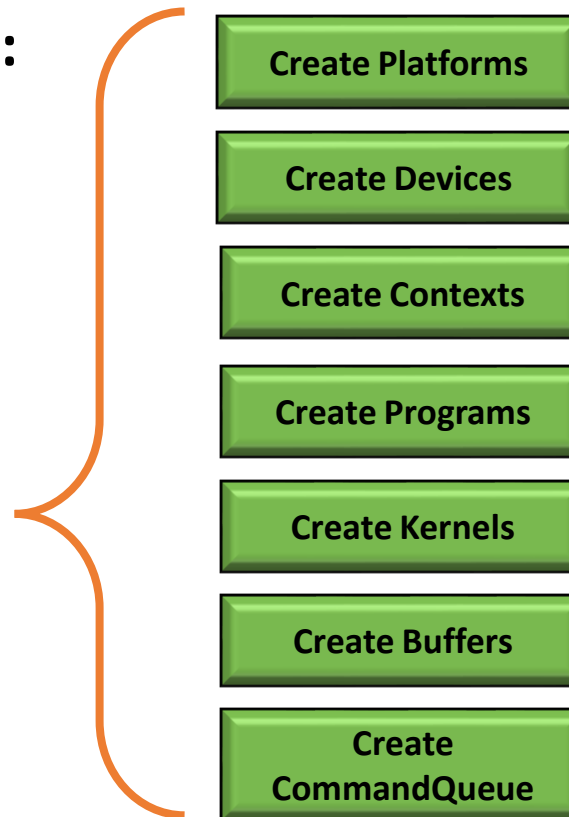
Hello World

- **Preparing two source code:**

- Hello_world.c
- Hello_world.cl

- **Program flow**

1. Read Hello_world.cl file
2. Create OpenCL structure
3. Send data to GPU
4. Send the task to GPU
5. Get the result from GPU
6. Done!



1. Read Hello_world.cl file

```
/* Create program from a file */
char *read_cl_file(const char* filename){
    FILE *program_handle;
    char *program_buffer;
    size_t program_size;

    /* Read program file and place content into buffer */
    program_handle = fopen(filename, "r");
    if(program_handle == NULL) {
        perror("Couldn't find the program file");
        exit(1);
    }
    fseek(program_handle, 0, SEEK_END);
    program_size = ftell(program_handle);
    rewind(program_handle);
    program_buffer = (char*)malloc(program_size + 1);
    program_buffer[program_size] = '\0';
    fread(program_buffer, sizeof(char), program_size, program_handle);
    fclose(program_handle);

    return program_buffer;
}
```

2. Create OpenCL structure (1/2)

```
int main() {  
  
    /* OpenCL data structures */  
    cl_platform_id platform;  
    cl_device_id device;  
    cl_context context;  
    cl_command_queue queue;  
    cl_program program;  
    cl_kernel kernel;  
    cl_int err;  
  
    /* Data and buffers */  
    char msg[16];  
    cl_mem msg_buffer;
```

2. Create OpenCL structure (2/2)

Create Platforms

Create Devices

Create Contexts

Create Programs

Create Kernels

Create Buffers

Create
CommandQueue

```
/* Create a platform */
err = clGetPlatformIDs(1, &platform, NULL);

/* Create a device */
err = clGetDeviceIDs(platform, CL_DEVICE_TYPE_ALL, 1, &device, NULL);

/* Create a context */
context = clCreateContext(NULL, 1, &device, NULL, NULL, &err);

/* Create a program */
char *programSource;
programSource = read_cl_file("Hello_world.cl");
program = clCreateProgramWithSource(context, 1, (const char**)&programSource, NULL, &err);

/* Compile a program */
clBuildProgram(program, 1, &device, NULL, NULL, NULL);

/* create a kernel */
kernel = clCreateKernel(program, "hello_kernel", &err);
```

3. Send data to GPU

```
/* Create a buffer to hold the output data */  
msg_buffer = clCreateBuffer(context, CL_MEM_WRITE_ONLY, sizeof(msg), NULL, &err);  
/* Create kernel argument */  
err = clSetKernelArg(kernel, 0, sizeof(cl_mem), &msg_buffer);  
/* Create a command queue */  
queue = clCreateCommandQueue(context, device, 0, &err);  
  
...  
...  
...
```

Create Platforms

Create Devices

Create Contexts

Create Programs

Create Kernels

Create Buffers

Create
CommandQueue

4. Send the task to GPU
5. Get the result from GPU
6. Done! (Release structure)

```
/* Enqueue kernel */
size_t globalWorkSize[1];
globalWorkSize[0]=2;
err = clEnqueueNDRangeKernel(queue, kernel, 1, NULL, globalWorkSize, NULL, 0, NULL, NULL);
/* Read and print the result */
err = clEnqueueReadBuffer(queue, msg_buffer, CL_TRUE, 0, sizeof(msg), &msg, 0, NULL, NULL);
printf("Kernel output: %s\n", msg);
/* Deallocate resources */
clReleaseMemObject(msg_buffer);
clReleaseKernel(kernel);
clReleaseCommandQueue(queue);
clReleaseProgram(program);
clReleaseContext(context);
return 0;
}
```


Hello_world.cl

```
__kernel
void hello_kernel(__global char *msg)
{
    *msg = (char)('H', 'e', 'l', 'l', 'o', ' ', ' ', 'k', 'e', 'r', 'n', 'e', 'l', '!', '!', '!', '\0');
    printf("Thread[%d]: Hello, welcome come to Opencl world.\n", get_global_id(0));
}
```

Next lecture will discuss

- **OpenCL kernel code**
- **Memory model**
- **Synchronization**

References

- Introduction to OpenCL.
Cliff Woolley, NVIDIA
http://www.cc.gatech.edu/~vetter/keeneland/tutorial-2011-04-14/06-intro_to_opengl.pdf
- CUDA C Programming Guide
<http://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html#axzz3g8riFfwE>
- CUDA Overview.
Cliff Woolley, NVIDIA
<http://www.cc.gatech.edu/~vetter/keeneland/tutorial-2011-04-14/02-cuda-overview.pdf>
- Introduction to CUDA
James Balfour, NVIDIA Research
http://mc.stanford.edu/cgi-bin/images/f/f7/Darve_cme343_cuda_1.pdf
- The OpenCL Programming Book - Free HTML version
Publisher: Fixstars
Author: Ryoji Tsuchiyama, Takashi Nakamura, Takuro Iizuka, Akihiro Asahara, Jeongdo Son, Satoshi Miki
- OpenCL in action
Author: Matthew Scarpino